BrandGuardAI

Pre-publication brand governance platform

Catch off-brand, inaccurate and unsupported content before it goes live. Explain every issue. Keep humans in control.

Event stage	Round 1: The Plan
Problem statement	Brand-governance platform (Inkloom AI)
Product name	**BrandGuardAI**
Team name	**BrandIQ**
Participant	**Ishika** (solo participant): research, design, development, testing and documentation
Education	B.Tech Computer Science and Engineering, 2nd year

1. Problem Definition

Content now comes from employees, interns, agencies, freelancers and AI tools. Many of them have never read the full brand guidelines, which are often long, outdated PDFs. Errors reach the public before anyone notices.

What goes wrong

- **Incorrect or outdated facts:** wrong prices, expired offers, old product names.
- **Unsupported promotional claims:** for example "No. 1" or "100% guaranteed".
- **Inconsistent tone and terminology:** the same product is called three different names.
- **Off-brand visuals and accessibility gaps:** wrong logo or colours, missing alt-text, low contrast.

Why today's process fails

- Review is manual: a manager compares each post with a guideline PDF by hand. It is slow and inconsistent.
- Feedback is scattered across chat, email and comments, so there is no single history of what changed and why.
- Reviewers are told *that* something is wrong, but not *which rule* it breaks or how to fix it.

Evidence that the problem is real

Published cases and data (details and sources in Section 3): Google's AI chatbot demo contained a factual error and Alphabet lost about US$100 billion in market value that day; a tribunal held Air Canada responsible for wrong information given by its chatbot; and 95.9% of one million home pages had detectable accessibility failures. Primary user research is planned and will be reported in Round 2.

2. Target Users and Setting

Setting: companies, agencies, startups and student clubs where about 5 to 50 people create content and a brand or marketing lead is responsible for consistency.

User	What they need	Current pain
Content Creator (employee, intern, freelancer)	Quick feedback on whether content is on-brand, with clear fixes and no need to read a 40-page guideline.	Unsure about rules. Gets vague corrections late, after rework is already done.
Brand / Marketing Manager	Consistent brand voice, fewer review rounds, visibility into who makes which mistakes.	Checks every item by hand. Becomes the bottleneck.
Authorised Reviewer (legal, compliance, brand owner)	Proof that each claim is backed by an approved source. An audit trail of decisions.	Cannot trace which source supports a claim. Approvals are not recorded.
Admin	Control of guidelines, approved facts, roles and confidential assets.	Guidelines live in scattered files and go out of date.

3. Research and Evidence

This evidence comes from published public cases and data. Primary user research is planned and will be reported in Round 2.

A. Desk research (completed)

Source	What it shows	Lesson for our design
Google Bard demo, Feb 2023	A promotional demo contained a wrong claim about the James Webb Space Telescope; Alphabet lost about US$100 billion in market value that day.	Claims must be checked against approved facts before publication.
Air Canada chatbot, Feb 2024	A tribunal held the airline responsible for a refund policy its chatbot described incorrectly, finding it did not take reasonable care to ensure accuracy.	Organisations are accountable for AI-generated content; human review and an audit trail are needed.
WebAIM Million 2026	95.9% of one million home pages had detectable WCAG 2 failures; low-contrast text on 83.9% and missing alt-text on 53.1%.	Accessibility mistakes are very common and many are machine-detectable, so automated checks add real value.

Sources: gpb.org (Bard), manatt.com (Air Canada), prioritypixels.co.uk (WebAIM).

B. Primary research plan (to be reported in Round 2)

- **Method:** online survey (12 questions) and 3 to 5 short interviews with student club media teams, interns, freelancers and marketing staff.

Survey theme	Example question	Why it matters for the product
How common the problem is	Have you seen off-brand, incorrect or misleading content published by your team?	Confirms the size of the problem.
Type of mistake	Was it tone, facts, claims, visuals or accessibility?	Decides which checks to build first.
Current process	How is content reviewed today and how long does it take?	Gives a baseline to measure time saved.
Guideline use	Are written guidelines available, and are they actually read?	Tests the need for guideline-clarity insights.
AI usage	Do you create content with AI tools?	Tests the need for AI-content checks.
Control preference	Should AI only suggest, or also edit automatically?	Validates the human-in-control design.

C. What the research tells us so far

- Wrong claims cause real damage, so every claim must trace to an approved source.
- Organisations stay responsible for AI content, so human approval is essential.
- Accessibility errors are common and detectable, so the MVP includes checks.

4. Proposed Solution and Value Proposition

BrandGuardAI is a web platform that reviews content before publication against an organisation's own approved guidelines, facts, terminology and assets. For every issue it explains **what is wrong, why, which guideline or approved source applies, and how to fix it** while keeping the creator's original purpose. Humans stay in control of every decision.

Value proposition

- **For creators:** instant, understandable feedback, so they fix issues themselves and learn the guidelines.
- **For managers:** fewer review rounds and one place to see status, comments and history.
- **For reviewers:** every claim is linked to an approved source, with a full audit trail.
- **For the organisation:** fewer public mistakes and data on which guidelines are unclear.

End-to-end workflow

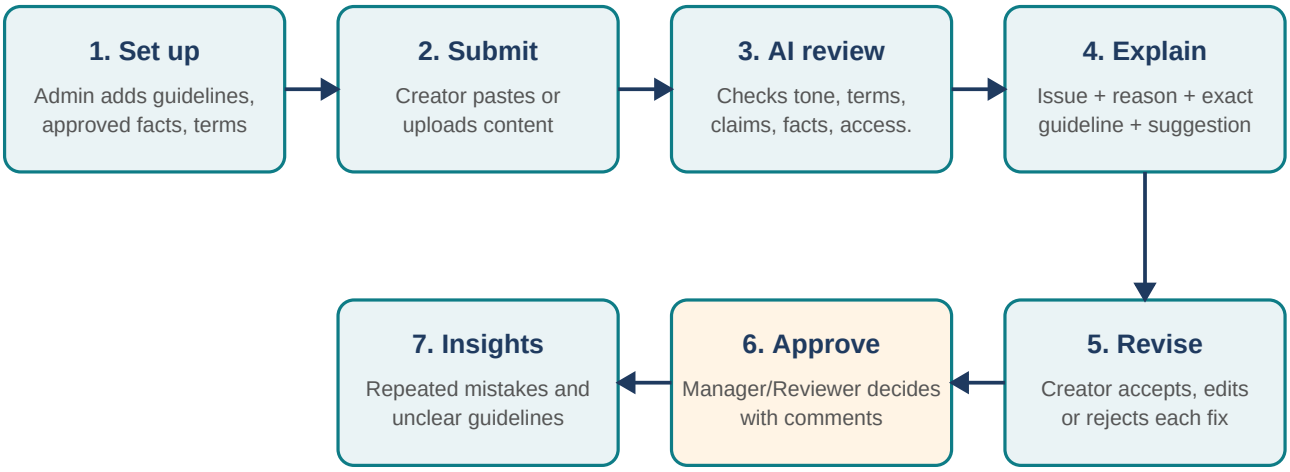

Human control: AI only suggests. Only an authorised reviewer can approve content.

Example: A creator writes "Our app is the #1 fastest in India." BrandGuardAI flags an *unsupported claim*, points to the guideline "Superlatives need an approved source", notes that no approved fact backs it, and suggests: "Our app is built for speed." The creator accepts, the manager approves, and the decision is recorded.

Product Preview: Wireframes

These wireframes show how the main screens are planned. The content in them is illustrative sample content, not real data.

Review screen

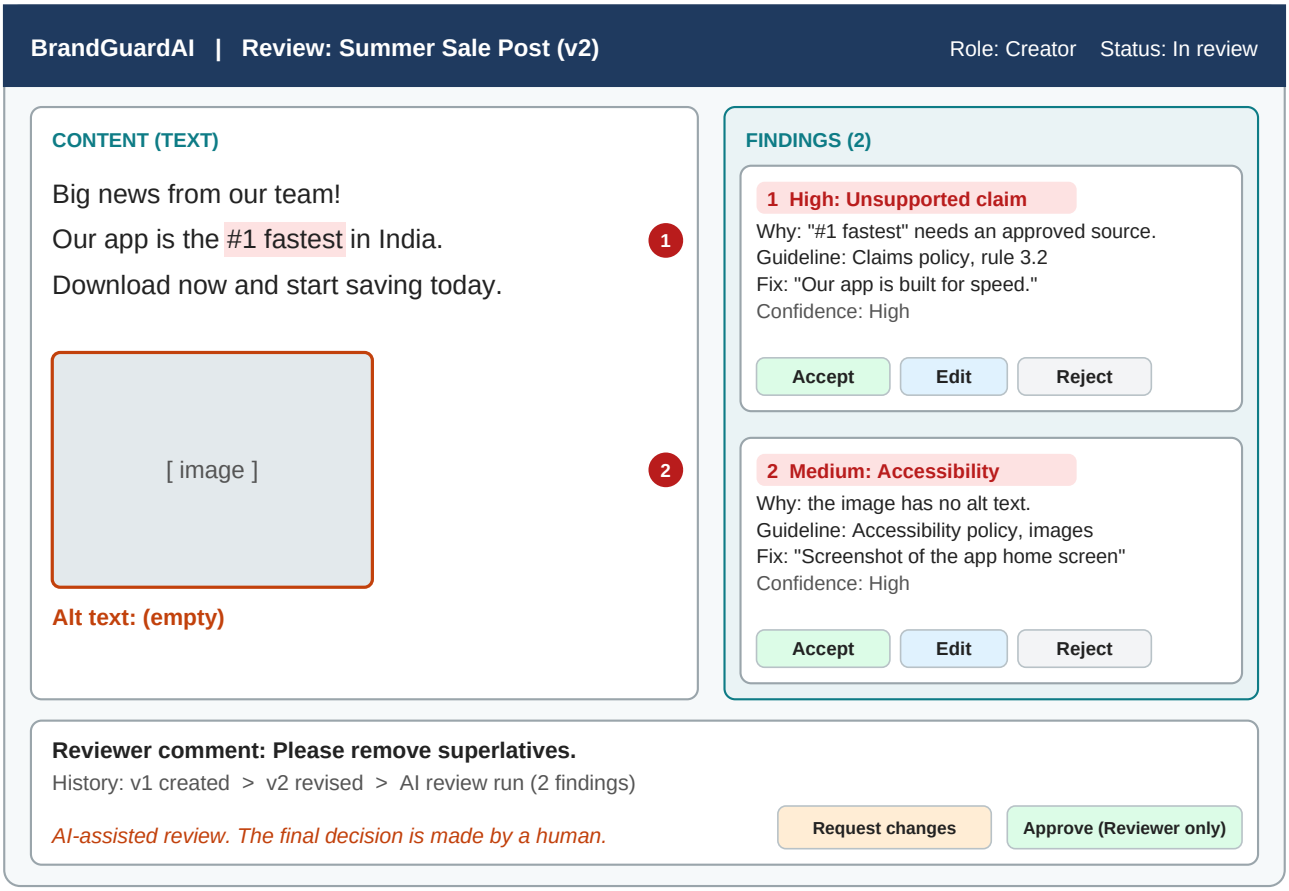

Each finding shows severity, the reason, the guideline involved, a suggested fix and a confidence level. The creator accepts, edits or rejects every suggestion. Only a Reviewer can approve.

Governance insights

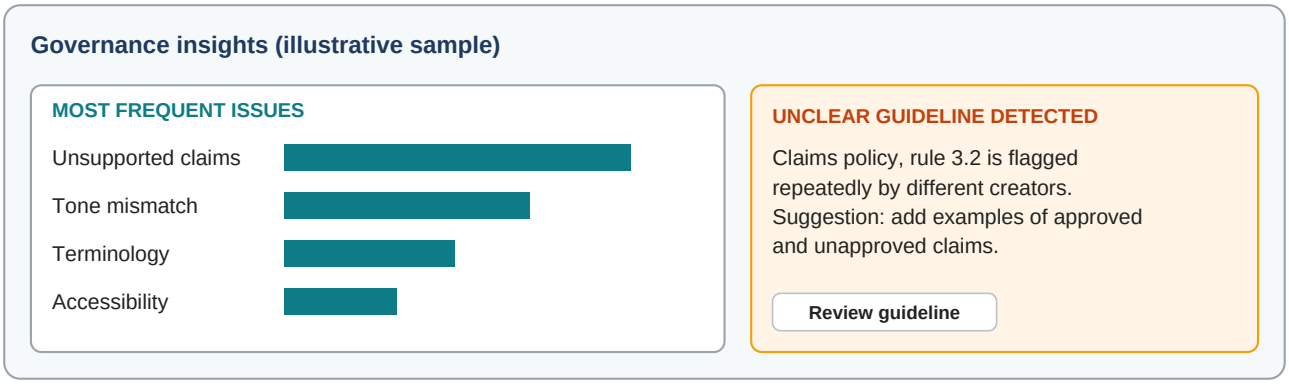

The dashboard shows the most frequent issues and points out guidelines that many creators get wrong, which suggests the guideline itself is unclear.

5. Planned Features

We deliberately limit the first version to a focused MVP so that the core workflow runs end to end and reliably.

Area	MVP (to be built)	Future roadmap
Brand knowledge	Upload guideline documents; define approved facts, terminology and banned/preferred words.	Versioned guidelines, asset library with usage rules.
Content types	Text: pasted text, TXT, DOCX, PDF.	Images (logo, colour checks), slides, video captions.
Detection	Tone, terminology, factual consistency, unsupported claims, basic accessibility (readability, alt-text presence).	Visual consistency, contrast checks, multilingual content.
Explanation	Each finding shows severity, reason, exact guideline or source quote, confidence and a suggested fix.	Learning from reviewer feedback.
Workflow	Roles (Creator, Reviewer, Admin), comments, approve/reject, revision history.	Notifications, Slack/Teams and CMS integration.
Insights	Dashboard of most frequent issues, repeat patterns, and guideline sections that cause confusion.	Trends per team, export reports.
Trust and safety	AI-use disclosure on every result, role-based access, audit log.	SSO, on-premise or private model option.

6. Originality of Approach

Many tools check grammar or tone. BrandGuardAI is different in four ways:

- **Source-cited findings.** Every flag names the exact guideline or approved fact involved, so reviewers can verify it in seconds.
- **Intent-preserving fixes.** We do not reject content. We suggest a correction that keeps the creator's purpose, and the creator chooses.
- **Guideline-clarity insights.** If many people repeat the same mistake on the same rule, the dashboard suggests the *rule* may be unclear. This improves the guidelines, not just the content.
- **Honest uncertainty.** The system shows confidence and escalates to a human when it cannot support a verdict, instead of producing confident but unreliable answers.

7. Technical Plan

Layer	Planned choice	Reason
Frontend	React (or HTML/CSS/JS)	Fast to build; clear review and dashboard screens.
Backend	Python (FastAPI)	Simple REST APIs; strong AI/document libraries.
AI	LLM API (for example Claude or GPT) with retrieval of guideline passages	Gives explainable results grounded in approved sources.
Database	SQLite for prototype, PostgreSQL when scaled	Stores users, content versions, findings, comments, audit log.
File parsing	PDF/DOCX text-extraction libraries	Import guidelines and documents.
Hosting	Render / Vercel (free tiers)	Judge-accessible deployed link.

How the review engine works

- **Step 1, Ingest:** guidelines and approved facts are split into small passages and stored with their source and section.
- **Step 2, Retrieve:** for each piece of content, the most relevant passages are found (retrieval-augmented generation).
- **Step 3, Check:** deterministic rules (banned words, required terms, missing alt-text) run in plain code. Subtler checks (tone, claim support) use the LLM, which is instructed to use *only* the retrieved passages.
- **Step 4, Structured output:** every finding is returned as: issue, severity, reason, source reference, suggested fix, confidence.
- **Step 5, Safeguard:** if no source supports a verdict, or confidence is low, the item is marked **"Needs human review"** instead of guessing.

Data sources and integrations

- Only the organisation's own uploaded guidelines, approved facts and assets. No scraping of third-party data.
- External services: one LLM API and a hosting provider. Both will be disclosed in the product.
- Integrations planned for later: Google Drive, Slack, CMS plug-ins.

Reliability measures

- Rules-first design so critical checks do not depend on the AI.
- Test set of sample posts with known issues, to measure how many issues are caught and how many flags are wrong.
- Reviewer "this flag was wrong" button, used to find weak spots.

8. Responsible Design

Area	Our approach
Human oversight	AI never edits, approves or publishes on its own. Suggestions are accepted or rejected by the creator; approval is only by an authorised reviewer.
Privacy and confidentiality	Each organisation's guidelines and assets are isolated and visible only to its members. Option to mask sensitive text before it is sent to an AI service. Minimum data is sent.
Security	Role-based access control, audit log of every action, encrypted storage and connections, API keys kept in environment secrets.
Transparency	Every AI-generated result is labelled. The product states which external AI service is used and what data is sent to it.
Explainability	Findings always include the reason and the source passage. Low-confidence results are marked for human review.
Fairness and errors	Reviewers can mark a flag as wrong; this feedback is tracked to find recurring false alarms.
Accessibility	The product's own interface will follow basic accessibility practice (keyboard use, contrast, labels).
Licences and IP	Open-source libraries used under their licences and disclosed. Only content users upload is processed.

9. Feasibility and Scope

As a solo participant, the scope is kept to the smallest set that proves the complete workflow, using existing libraries and one LLM API.

In scope for the build

- Text content review with explainable findings, suggestions and the approval workflow.
- Guideline and approved-fact import, roles, comments, revision history and a basic insights dashboard.

Out of scope initially (roadmap)

- Advanced image and video analysis, large-scale integrations, enterprise single sign-on.

Risks and how we handle them

Risk	Mitigation
AI gives wrong or invented findings	Ground answers in retrieved sources only, use rule-based checks for critical items, show confidence, require human decision.
Scope too large for one developer	MVP-first plan; roadmap features are only built if the core is stable.
Confidential data exposure	Per-organisation isolation, masking option, clear disclosure of external services.
API cost or downtime	Cache results, limit request size, fall back to rule-based checks.
Not enough user feedback	Run short user tests early and use the survey audience for testing.

10. Development Plan, Roles and Impact

Phase	Work	Outcome
Phase 1 (about 1 week)	Project setup, database design, guideline and fact import, document parsing.	Brand knowledge base working.
Phase 2 (about 1 week)	Review engine: rules, retrieval, LLM checks, structured explanations.	Text content gets explained findings.
Phase 3 (about 1 week)	Roles, comments, approval flow, revision history, insights dashboard.	Complete workflow end to end.
Phase 4 (about 1 week)	User testing, bug fixes, reliability tests, deployment, documentation and demo.	Stable, deployed MVP.

Phase durations are approximate and will be fitted to the official build window.

Role and responsibilities

Ishika (solo participant): all work, including user research, UI and frontend, backend and database, roles and approval workflow, AI review engine and testing, documentation and demo.

Expected impact

- Less time on manual review (to be measured during user testing).
- Fewer errors reach the public; consistent tone and terminology across teams.
- Clear audit trail for claims and approvals.
- Better guidelines over time through recurring-mistake insights.

Disclosure of AI assistance

As required by the event rules, we disclose that an AI assistant helped structure and draft this plan. All evidence, design decisions and final content were reviewed and confirmed by the participant. The planned product will use an external LLM API, which will be disclosed to users inside the product.